

Programming in Java

3. Multi-Threading 2



Java Multi-Threading

- As Java has evolved, there have been several versions of how Java implements multi-threading
 - The multi-threading shown in module 1 is the older version of multi-threading
 - Seen a lot of Java legacy code
- The more modern approach
 - Gets rid of a lot of the boiler plate code by not requiring us to write code to manage threads directly
 - Instead, the overhead of creating and running threads is moved to the Java class libraries
 - Instead we use threads through various interfaces
 - This also is more efficient from a resource utilization perspective
 - Java can optimize the number of threads running and ensure thread safety

The Resource Problem


- Threads are often used for task that are:
 - Short in duration
 - Called very frequently
- The problems with managing any kind of resource with these characteristics are
 - The amount of time spent creating and shutting down threads starts to become significant – the system starts to “thrash” trying to manage the threads
 - Too many threads can cause out of memory issues
- Managing threads at the low level we have been doing:
 - Requires writing a lot of boilerplate code
 - Is time consuming and error prone
- The solution is to delegate the actual creating and running of threads to the JRE

Pooling Resources

- Resource pool - collection of pre-created resources that are available on demand
 - This is a standard architectural strategy
 - Flyweight design pattern
- A thread pool contains a number of pre-created threads
 - Delay introduced by thread creating is eliminated, a thread is just selected from the pool
 - The thread is passed a Runnable object and then executes it
 - When the thread finishes executing, it is returned to the pool to be reused later
- Reduces thread life-cycle overhead and thrashing
- Allows resource limits to be set, like the maximum number of threads
- Allows programmers to concentrate on the executable code inside the run() method instead of the overhead of thread management
- This is the preferred way to create and use threads in modern Java programming

Executor Services

- The Executor interface
 - Creates a pool of threads and a queue to hold jobs
 - A runnable object to be executed in a thread is called a “job”
 - When a request comes in via a job queue, it is allocated to a thread which performs the task
 - When the task is finished, the thread is returned to the pool



Executor Services

- Java provides a concurrency library which supports a built-in Java thread pool
- Implemented as three interfaces
- An *Executor* interface that provides a replacement to the standard thread syntax.
 - *(new Thread(runnablecode)).start()* can be replaced by *e.execute(runnablecode)* where e is an instance of Executor
- *ExecutorService* interface extends the *Executor* interface to include a submit method for a *run()* method which returns a value
 - We will not be covering this interface in this class
- *ScheduledExecutorService* which adds methods to allow scheduling of threads

Executor Services

- Start by allocating a Thread pool using one of the constructors (factory method)
 - The following code implements a fixed size Executor service
 - The shutdown() message
 - *Stops the service from accepting new tasks*
 - *Shuts the service down when all the executing threads have exited*

```
public static void main(String[] args) {  
  
    // Creates a new Thread Pool with 3 executors  
    ExecutorService myPool = Executors.newFixedThreadPool(3);  
  
    // Shuts the pool down once all the threads have terminated  
    myPool.shutdown();  
  
}
```

Submit a Task

- Once the pool is started
 - Runnable tasks are submitted via the `execute()` method
 - `Execute` queues up the task, and when a thread is available, passes the task to the thread then executes the `start()` method on the thread

```
// Creates a new Thread Pool with 3 executors

ExecutorService myPool = Executors.newFixedThreadPool(3);
for (int i = 1; i < 5; i++) {
    myPool.execute(new MyTask("Task " + i));
}
// Shuts the pool down once all the threads have terminated
myPool.shutdown();
```


Customized Executor

- We can also create our own service with customized parameters
 - Core threads – the number of threads to start with
 - Max threads – the number of threads that can the executor service can scale up to
 - Keep alive – the amount of time to keep an executor running when idle
 - Time units – the time units used to measure the keep alive
 - BlockingQueue – the queue object to be used by the pool
- Executors are designed to be highly configurable

Customized Executor Submission

- Exactly the same as before
- The parameters are tuned for performance based on our requirements
- And based on performance history

```
public static void main(String[] args) {  
  
    int corePoolSize = 3;  
    int maxPoolSize = 5;  
    long keepAliveTime = 3000;  
    BlockingQueue<Runnable> pool = new ArrayBlockingQueue<Runnable>(100);  
  
    ExecutorService myPool = new ThreadPoolExecutor(  
        corePoolSize, maxPoolSize, keepAliveTime, TimeUnit.MILLISECONDS, pool);  
}
```

Lab 3-1

Executors



Futures

- Java programs often run tasks in separate threads.
 - In the examples we have seen so far, runnable tasks have been of the form `void run()`
 - But if we want to provide arguments and get a return value
- Basic issues to be solved are
 - How to start a task with parameters
 - How to determine when the task is finished
 - How we can get back the return value when the task is finished
- Futures are Java's way to handle asynchronous results.
 - Executors are designed to be highly configurable in order to handle Futures

Futures

- A `Future<T>` represents the result of an asynchronous computation.
 - Think of it like a placeholder for a value that's coming later.
 - The `T` is the type of data returned by the task
- For example
 - `Future<Integer> result = executor.submit(() -> 42);`
 - Submits a Lambda (runnable) to an executor that returns the value 42
 - When the task finishes, the value is stored in the result variable
 - You can think of the future as a box that will, at some point, hold a value
- You can:
 - Check if it's done → `future.isDone()`
 - Block and get the result → `future.get()`

Futures

- A `Future<T>` represents the result of an asynchronous computation.
 - Think of it like a placeholder for a value that's coming later.
 - The `T` is the type of data returned by the task
- For example
 - `Future<Integer> result = executor.submit(() -> 42);`
 - Submits a Lambda (runnable) to an executor that returns the value 42
 - When the task finishes, the value is stored in the result variable
 - You can think of the future as a box that will, at some point, hold a value
- You can:
 - Check if it's done → `future.isDone()`
 - Block and get the result → `future.get()`

Lab 3-2

Futures



Completable Futures

- Futures are limited
 - They do not handle exceptions
 - They do not “chain” meaning one task can call another task
 - Excellent for small individual tasks to be run in a thread
- `CompletableFuture` extends the capabilities of `Future` by allowing:
 - Non-blocking operations
 - Chaining of dependent tasks
 - Combining multiple futures
 - Handling results and exceptions fluently
 - Manual completion and cancellation
- Heavily influenced by reactive streams programming
 - Not covered in this course

Completable Futures

- Uses the ForkJoinPool
 - ForkJoinPool is a specialized thread pool designed for parallelism
 - Optimized for "divide-and-conquer" tasks, where large tasks are split (forked) into subtasks and then combined (joined).
- How It Works
 - Tasks are submitted to the pool.
 - Each thread in the pool has its own deque (double-ended queue).
 - When a thread finishes its work, it tries to steal a tasks waiting in other threads deques
 - Unlike an executor, each thread has it's own task buffer so there is no thread contention for a shared buffer

Completable Futures

- Downsides
 - Be careful overusing the common pool
 - It's shared, so heavy tasks may delay unrelated async operations.
- Avoid blocking operations in tasks submitted to the common pool.
 - It's optimized for CPU-bound tasks, not I/O.
 - Consider a custom thread pool for I/O-heavy or long-running tasks.

Non-Blocking Architecture

- Future.get() blocks the calling thread until the result is ready
 - CompletableFuture allows reacting to results asynchronously using callback methods
 - These are methods like .thenApply(), .thenAccept()
 - These take the result and then respond to it
 - Improves thread efficiency
 - No threads are sitting idle waiting for results.
- The callbacks do not block the code
 - The code continues to execute
 - At some point the task finished, then the callback is executed

Non-Blocking Example

- *CompletableFuture.supplyAsync()* -> *fetchFromDB()*
 - Submits *fetchFromDB()* to run asynchronously in the default *ForkJoinPool.commonPool()*
 - Returns immediately with a *CompletableFuture<T>* which is empty
 - The main thread continues execution without waiting.
 - *.thenAccept(result -> updatePage(result))* is a callback to be triggered when the async result is ready.
 - The code *updatePage(result)* will not run immediately, but will run after *fetchFromDB()* completes

```
CompletableFuture.supplyAsync(() -> fetchFromDB())  
    .thenAccept(result -> updatePage(result));
```

Chaining and Task Pipelines

- A task pipeline
 - Each task processes the result of the previous one.
- `CompletableFuture` enables this fluently using methods like:
 - `thenApply()` – transforms result
 - `thenAccept()` – consumes result
 - `thenRun()` – runs next action regardless of result

```
CompletableFuture.supplyAsync(() -> getUserId())  
    .thenCompose(id -> fetchProfile(id))  
    .thenApply(profile -> enrich(profile))  
    .thenAccept(this::sendToClient);
```

Composability

- Combine multiple independent or dependent tasks using:
 - *thenCombine()* – combines results of two independent tasks
 - *thenCompose()* – chains dependent async tasks
 - *allOf()* / *anyOf()* – coordinates many task
- Enables parallelism:
 - run tasks concurrently when they don't depend on each other.
- Enables complex workflows:
 - define multi-stage pipelines involving multiple results.

```
CompletableFuture<Integer> price = getPriceAsync();
CompletableFuture<Integer> tax = getTaxAsync();

price.thenCombine(tax, (p, t) -> p + t)
    .thenAccept(System.out::println);
```

Error Handling

- Can catch and respond to exceptions using:
 - exceptionally() – supply fallback
 - handle() – process result or exception
 - whenComplete() – log or audit outcomes
- Avoids unchecked exceptions crashing the app
- Allows graceful recovery (fallbacks, retries)
- Centralizes error flow logic

```
fetchData()
    .thenApply(data -> parse(data))
    .exceptionally(ex -> {
        logError(ex);
        return defaultData();
    });
```

Asynch

- Synchronous Methods
 - thenApply, thenAccept, etc.
 - Run the next step in the same thread that completed the previous stage.
- Asynchronous Methods
 - thenApplyAsync(), thenAcceptAsync()
 - These methods always run the next step in a different thread
- Notice in the examples
 - The tasks are all provided as Lambda functions

```
CompletableFuture<String> future = CompletableFuture.supplyAsync(() -> {  
    System.out.println("Task 1 in: " + Thread.currentThread().getName());  
    return "Hello";  
}).thenApply(result -> {  
    System.out.println("Task 2 in: " + Thread.currentThread().getName());  
    return result + " World";  
});
```

```
CompletableFuture<String> future = CompletableFuture.supplyAsync(() -> {  
    System.out.println("Task 1 in: " + Thread.currentThread().getName());  
    return "Hello";  
}).thenApplyAsync(result -> {  
    System.out.println("Task 2 in: " + Thread.currentThread().getName());  
    return result + " World";  
});
```


Lab 3-3

Completable Futures





Java™