

# Programming in Java

## 6. Java Streams



# Java Streams

- Functional programming support was implemented in Java to support stream processing
- A stream processes a collection of data objects
  - It takes input from a source of some kind without altering the source
  - All of the objects in a stream must be the same type
  - The data items move through a pipeline of transformations
  - A terminal operation ends the stream
- A stream in Java is not like a message queue
  - It can be helpful to think of it as a sequence of data objects
  - Conceptually, all the elements in a stream are processed at the same time at each step of the pipeline
  - A terminal operation either returns some collection or some single result or performs an operation (like saving to persistent storage) for each element in the stream
  - The actual pipeline is optimized before any processing takes place

# Why Streams

- One of the main drivers for using streams is to handle the requirements of big data
  - Data is often streaming from multiple sources
  - We want to process the items in the stream with some sort of transformation
  - We can't do this with OO or procedural programming
- If we use functional programming, then we have an elegant solution
  - We often need to write code that manages data streams for streaming data platforms
  - For example, Kafka has a streams capability that we program with Java streams
  - Data transformations on data streams in Spark is coded in Java, and other languages

## Some Notation

- Java has a `forEach(f)` method that applies to collections
  - It applies the the function `f` to each element of the list in turn
  - Sort of a functional programming version of a for loop.
- In the examples that follow
  - We use a short hand for a lambda function
  - Assume we have a collection or stream with an iterable interface
  - Then `forEach(System.out::println)`
  - `System.out::println` is a method reference.
  - Shorthand for `x -> System.out.println(x)` — a lambda that prints its input.
  - Common notation in streams.

# Initial Methods - Collections

- These take data source as input and return a stream
  - Any class implementing the Collection interface (e.g. List, Set) can produce a stream of its elements via the stream() method
  - The example to the right returns a sequential stream by default.
  - A List<String> can be turned into a Stream<String> for processing

```
List<String> items = Arrays.asList("apple", "banana", "cherry");  
  
// Create a sequential stream from the list  
Stream<String> stream = items.stream();  
  
// Use stream operations (e.g., convert to uppercase and print each)  
stream.map(String::toUpperCase)  
      .forEach(System.out::println);  
// Output: APPLE \n BANANA \n CHERRY (each on a new line)
```

# Initial Methods - Arrays

- There are two common ways to create a stream from an array
  - Using `Arrays.stream(...)`: This method takes an array and returns a sequential stream of its elements.
  - Using `Stream.of(...)`: This is a varargs method that can accept individual values or an array reference.

```
String[] fruits = { "apple", "banana", "cherry" };

// 1. Using Arrays.stream() on an array
Stream<String> fruitStream1 = Arrays.stream(fruits);
fruitStream1.forEach(System.out::println);
// Output: apple \n banana \n cherry

// 2. Using Stream.of() with explicit values
Stream<String> fruitStream2 = Stream.of("apple", "banana", "cherry");
fruitStream2.forEach(System.out::println);
// Output: apple \n banana \n cherry
```

# Initial Methods - Files

- You can create a stream from a file, typically to process lines of text.
  - The Files.lines(Path) method (in java.nio.file.Files) returns a Stream<String> where each element is a line in the file.
  - The example on the right reads a text file line by line by turning it into a stream of Strings

```
Path path = Paths.get("data.txt");
try (Stream<String> lines = Files.lines(path)) {
    lines.filter(line -> line.contains("ERROR"))
        .forEach(System.out::println);
} catch (IOException e) {
    e.printStackTrace();
}
```



# Infinite Streams

- Java streams can be infinite (unbounded).
  - The methods `Stream.iterate()` and `Stream.generate()` create streams that potentially never end
  - Typically used with a limit or a short-circuiting operation
  - Short circuit operation is one that will cause the stream to terminate, like when some test predicate fails

```
Stream<Double> randomNumbers = Stream.generate(Math::random);  
randomNumbers.limit(3).forEach(System.out::println);  
// Example output: 0.495... \n 0.365... \n 0.789... (3 random doubles)
```

```
Stream<Integer> powersOfTwo = Stream.iterate(1, n -> n * 2);  
powersOfTwo.limit(5).forEach(System.out::println);  
// Output: 1 \n 2 \n 4 \n 8 \n 16
```

```
Stream<Integer> powersOfTwo = Stream.iterate(1, n -> n * 2);  
powersOfTwo.limit(5).forEach(System.out::println);  
// Output: 1 \n 2 \n 4 \n 8 \n 16
```



# Parallel Streams

- A stream can be either sequential or parallel.
  - Parallel streams divide the workload across multiple threads, potentially speeding up processing on large data sets.
  - You can create a parallel stream in two ways:
    - *From a collection: use `parallelStream()` instead of `stream()` on a `Collection`. For example, `list.parallelStream()` produces a parallel stream*
    - *From an existing stream: call `stream.parallel()` to switch a sequential stream into parallel mode.*
- We will not cover this in this class

```
List<String> names = Arrays.asList("Alan", "Bob", "Cathy", "Doug");
names.parallelStream().forEach(System.out::println);
// Output (order may vary in parallel):
// Doug
// Cathy
// Alan
// Bob
```

# Pipeline Methods

- These take a stream as input
- Return a stream as output
- Large library of methods - some of these are:
  - **map**(function) – applies the function to each element of the stream
  - **filter**(predicate) – keeps the elements that match the predicate, discards the others
  - **sorted**() – sorts the stream
- Other pipeline methods are in the `java.util.streams` library

# Terminal Methods

- Terminal methods are methods that take an input from a stream and produce a final result
- Terminal methods mark the end of a stream – each stream can have only one terminal method
- Some terminal methods are:
  - **collect**(collection) – returns the result of the intermediate operations as a collection (e.g. list, array etc)
  - **forEach**(function) – applies the function to each element of the stream – does not produce an output stream
  - **reduce**(function) – uses function to collapse a stream into a single value

# Java Streams Simple Example

```
public static void main(String[] args) {  
  
    Function<Integer,Integer> square = x -> x * x;  
    // Source Collection  
    List<Integer> numbers = Arrays.asList(1,2,3,4,5,6,7,8,9,10);  
    List<Integer> squares =  
        numbers.stream()           // creates the stream object  
        .map(square)               // pipeline method applies the  
                                   // square function all the elements in the stream  
        .collect(Collectors.toList()); // Convert the stream to a list  
    System.out.println(numbers); // this as not been changed  
    System.out.println(squares);  
}
```

# Java Streams Simple Lambda Example

```
public static void main(String[] args) {  
  
    //Function<Integer,Integer> square = x -> x * x;  
    // Source Collection  
    List<Integer> numbers = Arrays.asList(1,2,3,4,5,6,7,8,9,10);  
    List<Integer> squares =  
        numbers.stream()           // creates the stream object  
        .map(x -> x * x)           // pipeline Lambda method applies the  
                                   // square function all the elements in the stream  
        .collect(Collectors.toList()); // Convert the stream to a list  
    System.out.println(numbers); // this as not been changed  
    System.out.println(squares);  
}
```

## Lazy Invocation

- Streams are not executed until a terminal method is encountered
- The stream is represented as a directed acyclic graph (DAG)
- This DAG can be optimized at compile time with a number of standard rewrite rules
  - The stream can only be optimized when the whole DAG is complete
  - And that happens when a terminal operation is encountered



# Lab 6-1

## Basic Java Streams





## Intermediate Methods

- Any stream method that returns a stream is an intermediate or pipeline method
- Some can be thought of as working on individual elements
  - Specifically, they can operate on an element without reference to other elements
  - Examples - `filter()`, `map()`
  - These operations can be parallelized
- Others need to examine the relationships between stream elements
  - Examples - `sorted()`, `distinct()`
- Intermediate operations should not have side effects
  - We violated this in some of the demos to see what was happening in a stream
  - The terminal methods are where any side effects should occur

## Intermediate Methods

- `map(f)`
  - Applies a monadic function `f` to each element in the stream and returns a transformed element
- `filter(p)`
  - Applies the predicate `p` to each element in the stream, if the result is false, the element is removed from the stream
- `peek(f)`
  - Used for debugging, it executes `f` on each element, when you want to see the elements as they flow past a certain point in a pipeline
- `distinct()`
  - removes duplicates based on the defined equality operator for the stream elements
- `sorted()`
  - sorts the elements based on the defined comparison operator for the stream elements

## Intermediate Methods

- `skip(n)`
  - Omits the first `n` elements of a stream
- `limit(n)`
  - Truncates the stream after the `n` elements
- `flatMap(stream)`
  - removes levels of structure to flatten a stream
  - For example, a list of lists of integers has two levels of structure
    - `[ [2, 3, 5], [7, 11, 13], [17, 19, 23] ]`
  - Flattening the list removes the nested structure
    - `[ 2, 3, 5, 7, 11, 13, 17, 19, 23 ]`

# Lab 6-2

## Pipeline Methods



# Terminal Operations

- The three basic types of terminal operations are
  - **Reducers** - returns a single value representing a computation on the stream - the stream is reduced to a single value
  - **Collectors** - returns some sort of collection
  - **Operators** - performs an operation on each element of the stream and returns void
- Terminal methods are terminal because they do not return a stream
  - They represent the end of the stream

# Reducers

- There are a number of standard reducers
- `count()`
  - Returns the number of elements in the stream
- `min(comparator)`, `max(comparator)`
  - These return Optionals which are like futures to account for the cases where the value may or may not be returned
  - If `isPresent()` is true meaning that the value exists, it can be retrieved with the `get()` method
  - The comparator is the predicate used to determine how to order the elements
- `anyMatch(p)`, `allMatch(p)`, `noneMatch(p)`
  - Returns a Boolean if the predicate `p` is
    - *true for any one of the elements in the stream*
    - *true for all the elements in the stream*
    - *true for none of the elements in the stream*

# Reducers

- `reduce(accumulator, operator)`
  - The accumulator is the last value computed (ie. from the previous element)
  - The operator is a function applied to combine the accumulator with the current element
  - Like summing an array in a loop
    - *The accumulator is the running total*
    - *The operator is adding the current element to the running total*



# Collector

- The the Collectors class has a number of methods that return collections of various types
  - Eg. toList(), toMap(), toSet()
- There are other sorts of collectors that, for example:
  - Combine the stream into a single String
  - Do reduction type operations as well
  - In fact, reducers can be thought of as special cases of collectors

# Lab 6-3

## Terminal Methods







Java™